

### **Unit III : Address Spaces**


- Early Systems, Multiprogramming and Time Sharing, The Address Space,
- Memory API: Types of Memory, The malloc() Call, The free() Call, Common Errors, Underlying OS Support,
- Segmentation, Fine-grained vs. Coarse-grained Segmentation, Free-Space Management,
- Paging, A Memory Trace, Faster Translations, (TLBs), TLB Basic Algorithm, Example: Accessing An Array, Who Handles The TLB Miss?, TLB Issue: Context Switches, Replacement Policy,
- Hybrid Approach: Paging and Segments, Beyond Physical Memory: Mechanisms, Swap Space, The Present Bit, The Page Fault, What If Memory Is Full?, Page Fault Control Flow, When Replacements Really Occur, The Linux Virtual Memory System.

# Why Virtualize Memory ?

- Because real view of memory is messy.
- Earlier, memory had only code of one running process (and OS code)
- Now, multiple active processes timeshare CPU.
  - Many processes must be in memory
  - Non-contiguous too
- Need to hide this complexity from user

# Multiprogramming and Time Sharing

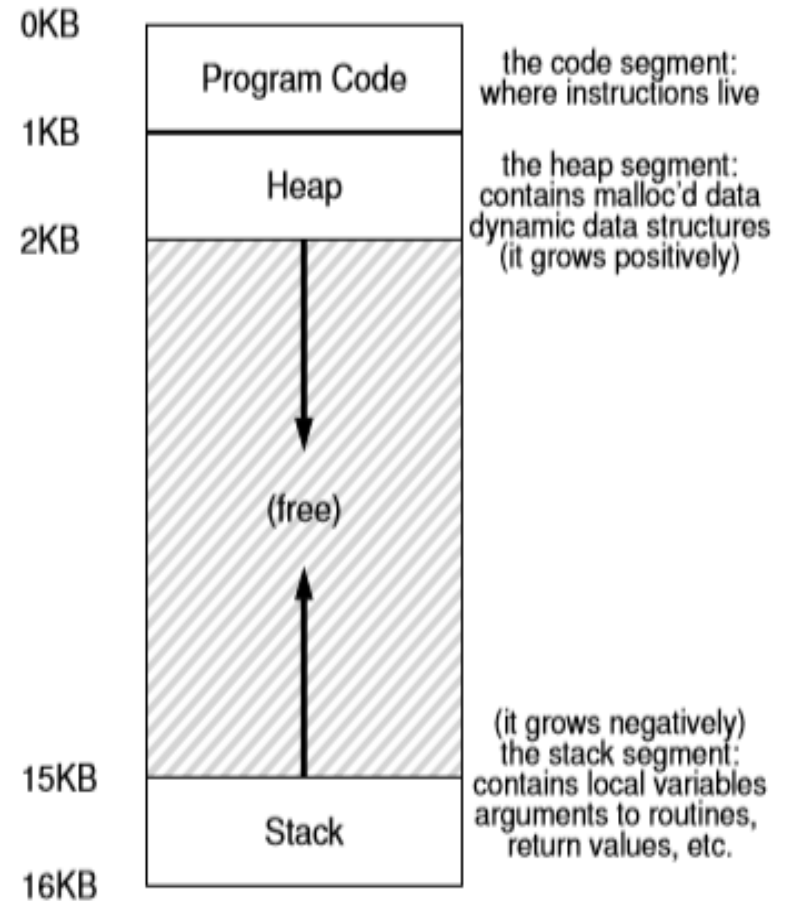
- ▣ **Load multiple processes** in memory.
  - ◆ Execute one for a short while.
  - ◆ Switch processes between them in memory.
  - ◆ Increase utilization and efficiency.



**Three Processes: Sharing Memory**

# Abstraction of Virtual Address Space

- OS creates an **abstraction** of physical memory, called as Address Space.
- Virtual Address Space: Every process assume it has access to large space of memory from address 0 to MAX
  - The address space contains all about a running process.
  - That is consist of program code, heap, stack etc.
- CPU issues load and store to virtual addresses



An Example Address Space

# Address Space(Cont.)

## □ Program


- ◆ Where instructions and static data live

## □ Heap

- ◆ Dynamically allocate memory.
  - `malloc/free` in C language
  - `new/delete` in object-oriented language


## □ Stack

- ◆ Store return addresses or values.
- ◆ Contain local variables arguments to routines.



Address Space

- Observe how processes A, B and C are loaded into memory
- When the OS does this, we say the OS is virtualizing memory, because the running program thinks it is loaded into memory at a particular address (say 0) and has a potentially very large address space (say 32-bits or 64-bits); the reality is quite different.




**Three Processes: Sharing Memory**

# Goals of VM system

- Transparency:
  - The program shouldn't be aware of the fact that memory is virtualized.
- Efficiency:
  - The OS should strive to make the virtualization as efficient as possible, both in terms of time and space.
  - In implementing time-efficient virtualization, the OS will have to rely on hardware support, including hardware features such as TLBs
- Protection:
  - Protection enables us to deliver the property of isolation among processes.

# Contiguous Memory

- Total size of memory is 10 KB
  - P1 requires 2 KB.
  - P2 requires 5 KB




- We can not store P2 in contiguous location,
- **Space is available but can no be allotted, Called as External Fragmentation.**
- **Advantages: Address translation and access is easy (only need to remember base address)**
- **Disadvantage: External Fragmentation**



# Non-Contiguous Allocation

- Total size of memory is 10 KB
  - P1 requires 2 KB.
  - P2 requires 5 KB (provides link)




We can not access the data element directly

Advantages: No external fragmentation.

# Contiguous : Fixed Partitioning


- Contiguous Technique can be divided into:
  - Fixed (or static) partitioning
  - Variable (or dynamic) partitioning
- Fixed Size Partitioning:
  - This is the oldest and simplest technique used to put more than one processes in the main memory.
  - In this partitioning, number of partitions (non-overlapping) in RAM are **fixed but size** of each partition may or **may not be same**



# Contiguous : Variable-sized Partitioning

- The size of partition will be equal to incoming process.
- The partition size varies according to the need of the process so that the internal fragmentation can be avoided to ensure efficient utilisation of RAM.
- Number of partitions in RAM is not fixed and depends on the number of incoming process and Main Memory's size.
- No internal fragmentation but external fragmentation.
  - For example, suppose in above example- process P1(2MB) and process P3(1MB) completed their execution. Hence two spaces are left i.e. 2MB and 1MB.
  - Let's suppose process P5 of size 3MB comes

## Dynamic partitioning



Block size = 2 MB

Block size = 7 MB

Block size = 1 MB

Block size = 5 MB

Partition size = process size  
So, no internal Fragmentation

- For both the schemes(fixed and variable size partitioning) , the operating system must keep list of each memory location noting which are free and which are busy.
- Then as new jobs come into the system, the free partitions must be allocated.
- These partitions may be allocated by 4 ways:
  - First Fit Memory Allocation
  - Best Fit Memory Allocation
  - Worst Fit Memory Allocation

# Memory API

# Memory API

- Memory allocation interfaces in UNIX system.
- In UNIX/C programs, understanding how to allocate and manage memory is critical in building robust and reliable software.
- What interfaces are commonly used?
- What mistakes should be avoided?

# Types of Memory

- In C program there are two types of memory allocated
- Stack: Allocation and deallocation is managed implicitly by the compiler.
  - Sometimes, called as automatic memory.
  - When you return from the function, compiler deallocates the memory

```
void func() {  
    int x; // declares an integer on the stack  
    ...  
}
```

- For long-lived memory:
  - Malloc: Allocation and deallocation is managed by programmer

```
void func() {  
    int *x = (int *) malloc(sizeof(int));  
    ...  
}
```

Compiler knows to make room for pointer to an integer when it sees (int \*)  
When it request malloc(), it requests space for an integer on the heap

# Memory API: malloc()

```
#include <stdlib.h>

void* malloc(size_t size)
```

- ▣ Allocate a memory region on the heap.
  - ◆ Argument
    - `size_t size` : size of the memory block(in bytes)
    - `size_t` is an unsigned integer type (defined in C standard).
  - ◆ Return
    - Success : a void type pointer to the memory block allocated by `malloc`
    - Fail : a null pointer



# Sizeof()

- ▣ Routines and macros are utilized for `size` in `malloc` instead typing in a number directly.
- ▣ Two types of results of `sizeof` with variables
  - ◆ The actual size of `'x'` is known at run-time.

```
int *x = malloc(10 * sizeof(int));  
printf("%d\n", sizeof(x));
```

4

- ◆ The actual size of `'x'` is known at compile-time.

```
int x[10];  
printf("%d\n", sizeof(x));
```

40


# Memory API: free()

```
#include <stdlib.h>

void free(void* ptr)
```


- ▣ Free a memory region allocated by a call to `malloc`.
  - ◆ Argument
    - `void *ptr`: a pointer to a memory block allocated with `malloc`
  - ◆ Return
    - none

# Memory Allocating




-----> pointer

```
int *pi; // local variable
```




```
pi = (int *) malloc(sizeof(int) * 4);
```

# Memory Freeing



```
free(pi);
```



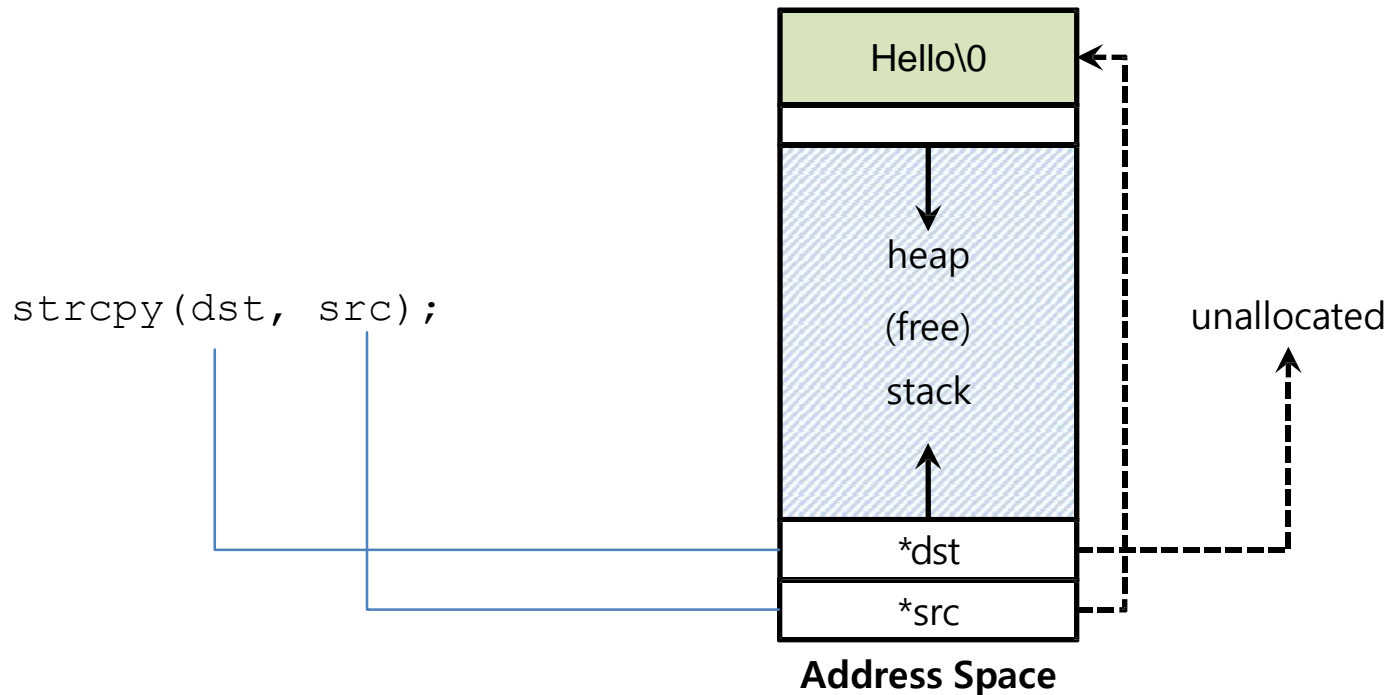
# Common errors (that arise in the use of malloc() and free())

- **Forgetting To Allocate Memory**
- **Not Allocating Enough Memory**
- **Forgetting to Initialize**
- **Forgetting To Free Memory/Memory Leak**
- **Freeing Memory Before You Are Done With It/Dangling Pointer**
- **Freeing Memory Repeatedly/Double free**
- **Calling free() Incorrectly/Invalid Free**

# Forgetting To Allocate Memory

## ■ Incorrect code :segmentation fault


```
char *src = "hello"; //character string constant
char *dst;           //unallocated
strcpy(dst, src);    //segfault and die
```



# Forgetting To Allocate Memory(Cont.)

## ▣ Correct code

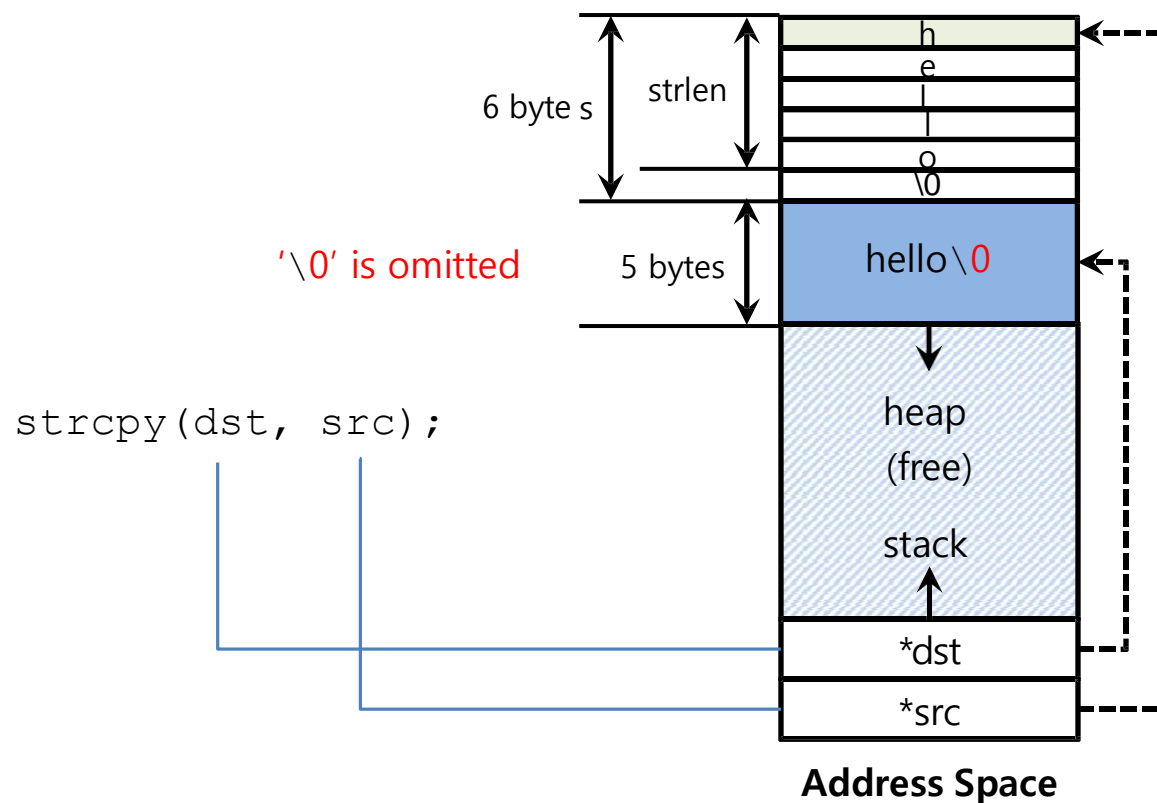
```
char *src = "hello"; //character string constant
char *dst = (char *)malloc(strlen(src) + 1 ); // allocated
strcpy(dst, src);    //work properly
```



# Not Allocating Enough Memory

- Incorrect code, but work properly

```
char *src = "hello"; //character string constant
char *dst = (char *)malloc(strlen(src)); // too
small strcpy(dst, src); //work properly
```







# Forgetting to Initialize allocated memory

- Encounter an uninitialized read

```
int *x = (int *)malloc(sizeof(int)); // allocated
printf("*x = %d\n", *x); // uninitialized memory
access
```



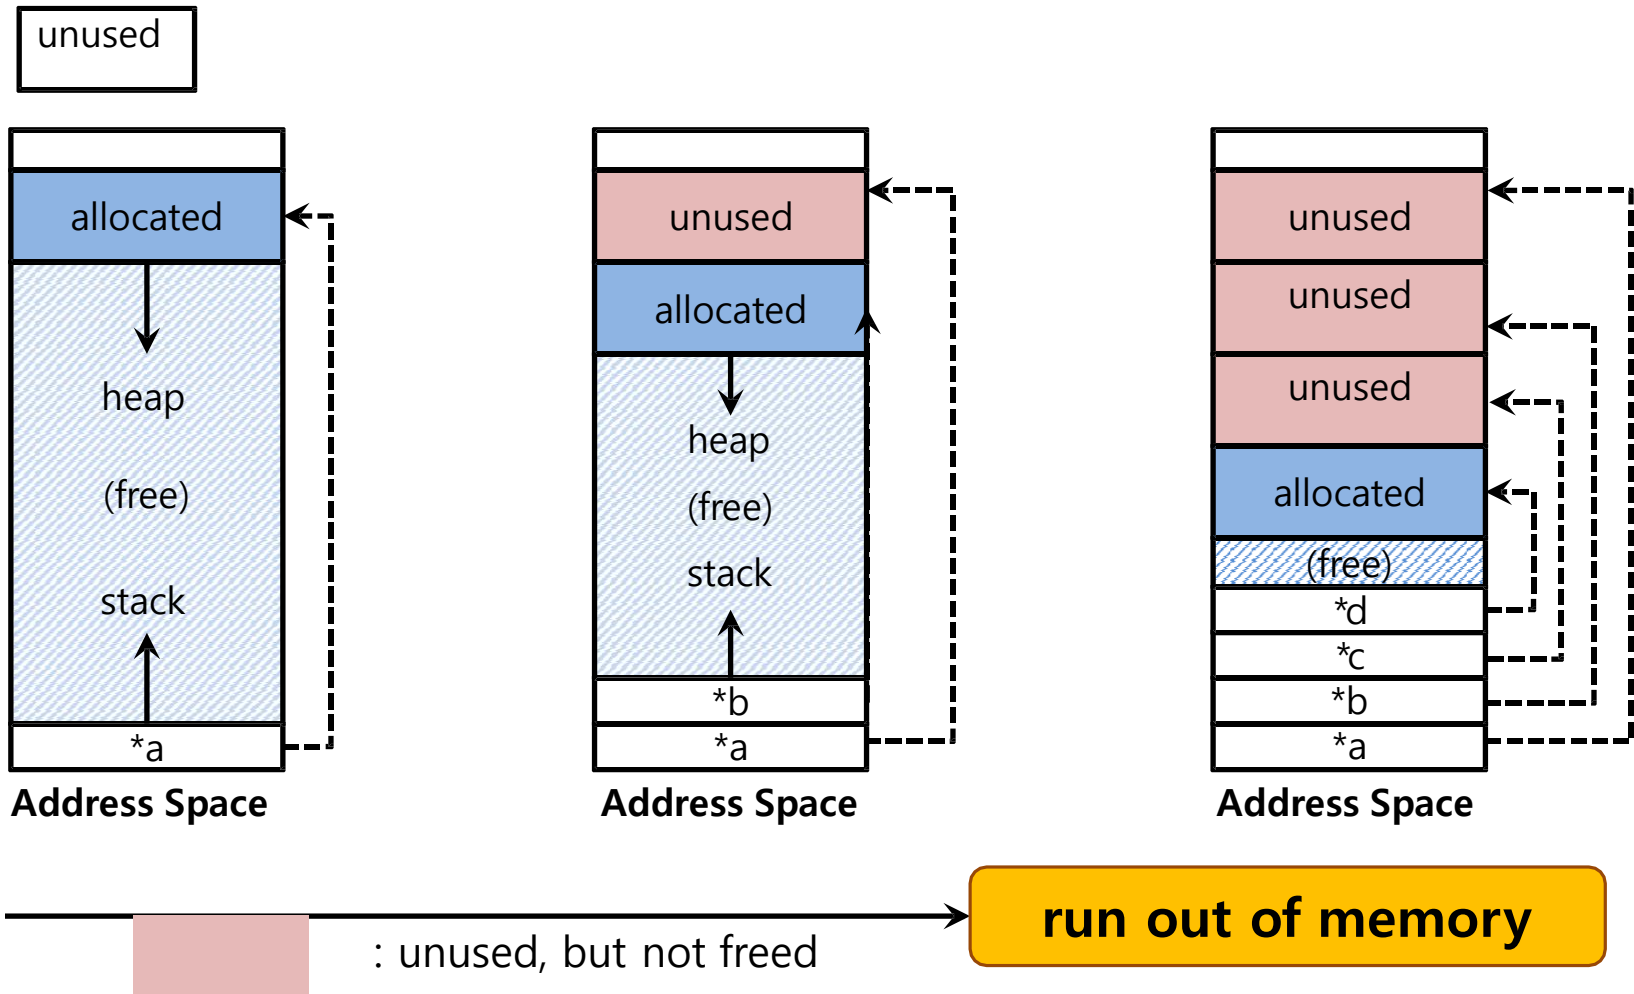
Address Space



Address Space

# Memory Leak




- ▣ A program runs out of memory and eventually dies.
- ▣ It occurs when you forget to free memory




- In long running applications or systems (such as OS) it is a huge problem.
- Slowly leaking memory leads one to run out of memory, because of restart is required.
- So, whenever we have done with chunk of memory it should be freed

# Dangling Pointer

- Freeing memory before it is finished using
  - A program accesses to memory with an invalid pointer




`free(b)`



# Double Free

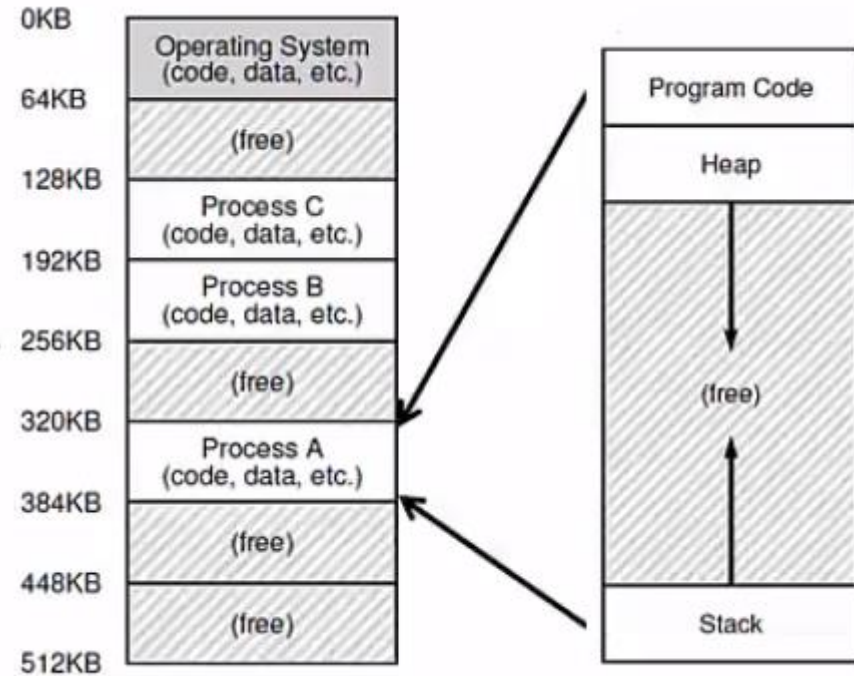
- Free memory that was freed already.

```
int *x = (int *)malloc(sizeof(int)); // allocated
free(x); // free memory
free(x); // free repeatedly
```



# How is actual memory reached?

- When program is running: CPU fetches load and store instructions
- All of these accesses happen to virtual addresses.
- But in real life, program is not stored from 0 to MAX
- When you access a certain address, somebody translated from VA to PA.
- CPU issues load/stores to VA but memory hardware accesses PA
  - OS allocates memory and tracks location of processes.
- Translation done by memory hardware called memory management Unit(MMU)



# Memory Allocation System Call

- Malloc implemented by c library
  - Algorithms for efficient memory allocation and free space management
- To grow heap, c library uses the brk/sbrk system call
- C library will take care of it.
- A program can also allocate a page sized memory using mmap() system call
  - Gets anonymous page from OS.
  - Gives page in the memory image of the process