

# Concurrency: An Introduction

---

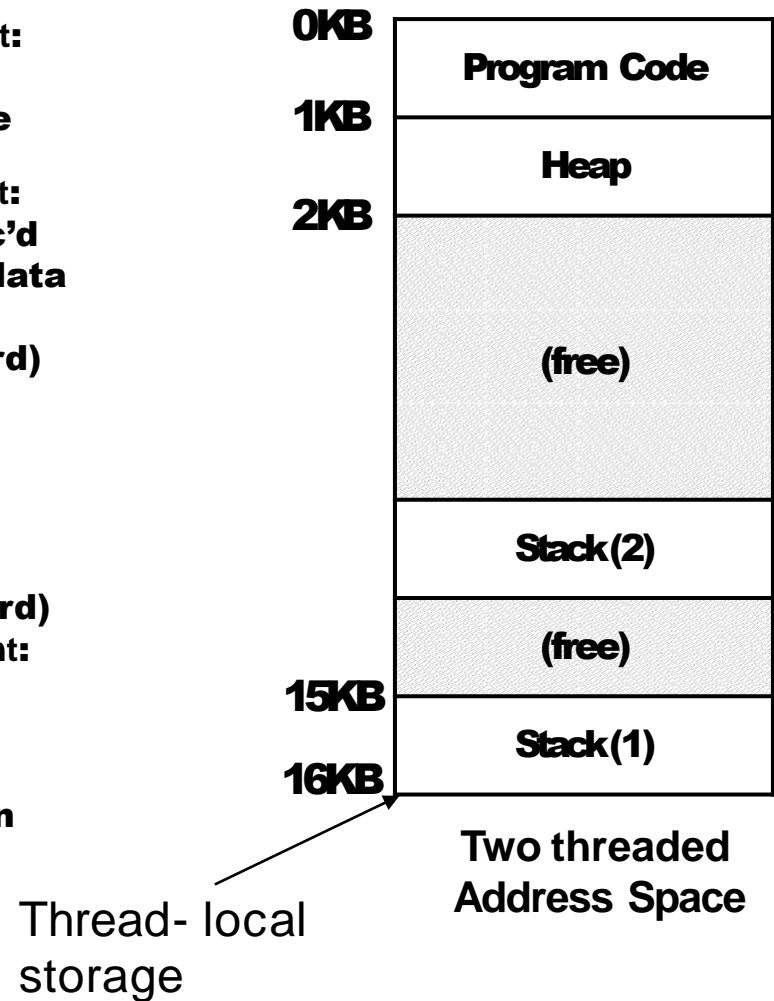
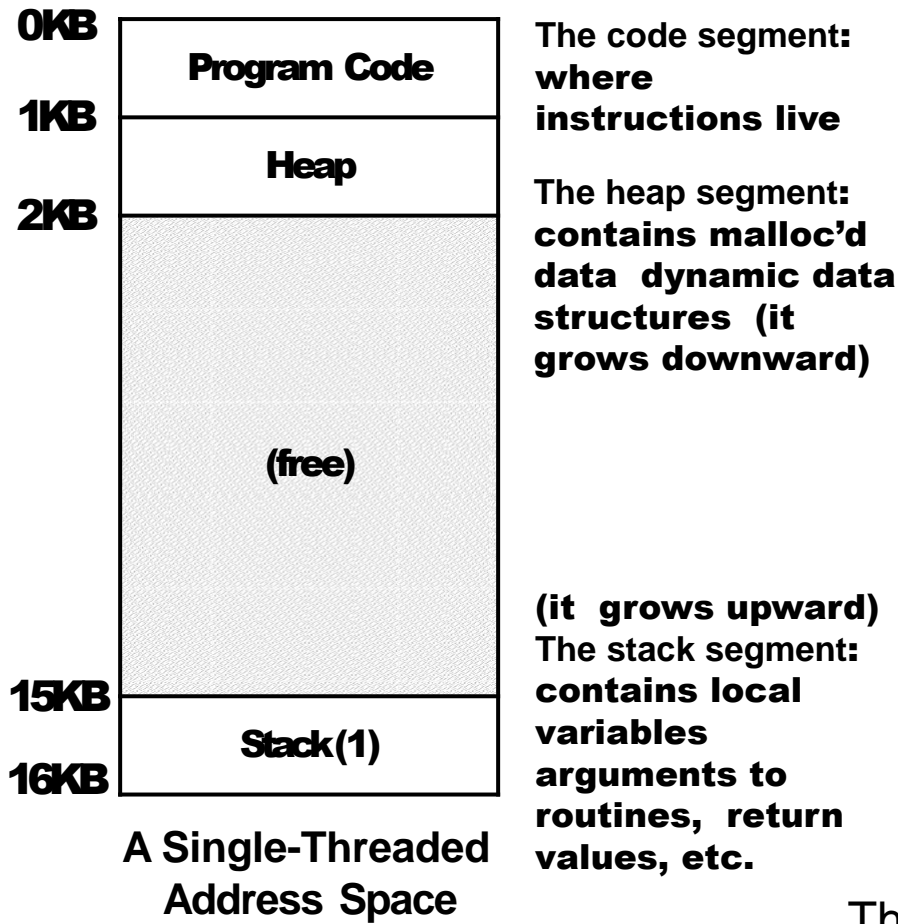
- ▣ A new abstraction for a single running process
- ▣ Multi-threaded program
  - ◆ A multi-threaded program has more than one point of execution.
  - ◆ Multiple PCs (Program Counter)
  - ◆ They **share** the same **address space**.

# Context switch between threads

- ▣ **Each thread has its own program counter and set of registers.**
  - ◆ **One or more thread control blocks(TCBs) are needed to store the state of each thread.**
  - ◆ **All of them within a common PCB**
  
- ▣ **When switching from running one (T1) to running the other (T2),**
  - ◆ **The register state of T1 be saved.**
  - ◆ **The register state of T2 restored.**
  - ◆ **The **address space remains** the same.**

# The stack of the relevant thread

- There will be **one stack per thread**.



# Why threads?

## ▣ Performance

- ◆ **Parallelism is the only way to use translate multiple cores into performance**
- ◆ **Parallelization: from single-threaded programs to multi-threaded**

## ▣ Convenience

- ◆ **Way to overlap I/O with useful work: approach of server-base applications such as web-servers, DBMS, etc..**

## ▣ Why threads and not processes?

- ◆ **In threads is much easier to share data**
- ◆ **Less pressure over the memory**
- ◆ **Processes when the task are separated with little (to none) sharing**

```
1  #include <stdio.h>
2  #include <assert.h>
3  #include <pthread.h>
4
5  void *mythread(void *arg) {
6      printf("%s\n", (char *) arg);
7      return NULL;
8  }
9
10 int
11 main(int argc, char *argv[]) {
12     pthread_t p1, p2;
13     int rc;
14     printf("main: begin\n");
15     rc = pthread_create(&p1, NULL, mythread, "A"); assert(rc == 0);
16     rc = pthread_create(&p2, NULL, mythread, "B"); assert(rc == 0);
17     // join waits for the threads to finish
18     rc = pthread_join(p1, NULL); assert(rc == 0);
19     rc = pthread_join(p2, NULL); assert(rc == 0);
20     printf("main: end\n");
21     return 0;
22 }
```

**Figure 26.2: Simple Thread Creation Code (t0.c)**

# Possible outcomes

main	Thread 1	Thread2
starts running		
prints "main: begin"		
creates Thread 1		
creates Thread 2		
waits for T1		
	runs	
	prints "A"	
	returns	
waits for T2		
		runs
		prints "B"
		returns
prints "main: end"		

main	Thread 1	Thread2
starts running		
prints "main: begin"		
creates Thread 1		
		runs
		prints "A"
		returns
	creates Thread 2	
		runs
		prints "B"
		returns
waits for T1		
<i>returns immediately; T1 is done</i>		
waits for T2		
<i>returns immediately; T2 is done</i>		
prints "main: end"		

main	Thread 1	Thread2
starts running		
prints "main: begin"		
creates Thread 1		
creates Thread 2		
		runs
		prints "B"
		returns
waits for T1		
	runs	
	prints "A"	
	returns	
waits for T2		
<i>returns immediately; T2 is done</i>		
prints "main: end"		

```

1          #include <stdio.h>
2          #include <pthread.h>
3          #include "mythreads.h"
4
5          static volatile int counter = 0;
6
7          //
8          // mythread()
9          //
10         // Simply adds 1 to counter repeatedly, in a loop
11         // No, this is not how you would add 10,000,000 to
12         // a counter, but it shows the problem nicely.
13         //
14         void *
15         mythread(void *arg)
16         {
17             printf("%s: begin\n", (char *) arg);
18             int i;
19             for (i = 0; i < 1e7; i++) {
20                 counter = counter + 1;
21             }
22             printf("%s: done\n", (char *) arg);
23             return NULL;
24         }
25
26         //
27         // main()
28         //
29         // Just launches two threads (pthread_create)
30         // and then waits for them (pthread_join)
31         //
32         int
33         main(int argc, char *argv[])
34         {
35             pthread_t p1, p2;
36             printf("main: begin (counter = %d)\n", counter);
37             Pthread_create(&p1, NULL, mythread, "A");
38             Pthread_create(&p2, NULL, mythread, "B");
39
40             // join waits for the threads to finish
41             Pthread_join(p1, NULL);
42             Pthread_join(p2, NULL);
43             printf("main: done with both (counter = %d)\n", counter);
44             return 0;
45         }

```



# Possible outcomes

```
prompt> gcc -o main main.c -Wall -pthread
prompt> ./main
main: begin (counter = 0)
A: begin
B: begin
A: done
B: done
main: done with both (counter = 20000000)
```

```
prompt> ./main
main: begin (counter = 0)
A: begin
B: begin
A: done
B: done
main: done with both (counter = 19345221)
```

```
prompt> ./main
main: begin (counter = 0)
A: begin
B: begin
A: done
B: done
main: done with both (counter = 19221041)
```

# The heart of the problem:: Uncontrolled Scheduling

## □ Example with two threads

- ♦ **counter = counter + 1 (default is 50)**
- ♦ **We expect the result is 52.**

However,  
OS

Thread1

Thread2

(after instruction)  
PC    %eax    counter

before critical section

mov 0x8049a1c, %eax

add \$0x1, %eax

100    0    50

105    50    50

108    51    50

interrupt

save T1's state (TCB)

restore T2's state (TCB)

mov 0x8049a1c, %eax

add \$0x1, %eax

mov %eax, 0x8049a1c

100    0    50

105    50    50

108    51    50

113    51    51

interrupt

save T2's state

restore T1's state

mov %eax, 0x8049a1c

108    51    50

113    51    **51**

- **Do the read and modification of the memory in a single step**
  - ♦ i.e. “all or nothing”!
- **How to handle complex data? (v.gr. a b-tree)**
  - ♦ Use some atomic hardware support (called **synchronization primitives**) to construct OS support
- **A piece of code that **accesses a shared variable** and must not be concurrently executed by more than one thread.**
  - ♦ Multiple threads executing critical section can result in a race condition.
  - ♦ Need to support atomicity for critical sections (mutual exclusion)

- **Ensure that any such critical section executes as if it were a single atomic instruction (execute a series of instructions atomically).**

```
1  lock_t mutex;  
2  . . .  
3  lock(&mutex);  
4  balance = balance + 1;  
5  unlock(&mutex);
```

**Critical section**

# One more problem: Waiting for another/s

- ▣ **Some times the thread interaction is wait for another thread**
  - ◆ **V.gr. When a thread should wait to another that had issued a I/O**
  - ◆ **Need to be slept until the other thread receives the I/O end**
- ▣ **Some times the action of multiple threads should be synchronous**
  - ◆ **V.gr. Many threads are performing in parallel a iteration in a numerical problem**
  - ◆ **All threads should start the next iteration at once (barrier)**
- ▣ **This sleeping/waking cycle will be controlled by condition variables**

# Interlude: Thread API

---

# Thread Creation

## ▣ How to create and control threads?

```
#include <pthread.h>

int
pthread_create(      pthread_t*      thread,
                    const pthread_attr_t* attr,
                    void*             (*start_routine) (void*),
                    void*             arg);
```

- ◆ `thread`: Used to interact with this thread (OUT).
- ◆ `attr`: Used to specify any attributes this thread might have.
  - Stack size, Scheduling priority, ... (IN)
- ◆ `start_routine`: the function this thread start running in (IN)
- ◆ `arg`: the argument to be passed to the function (`start_routine`) (IN/OUT)
  - *a void pointer* allows us to pass in *any type of* argument.
- ◆ Returns 0 if went good (a error code otherwise: EAGAIN, EINVAL, EPERM)

# Thread Creation (Cont.)

- ▣ If `start_routine` instead required another type argument, the declaration would look like this (example):

- ◆ An integer argument:

```
int
pthread_create(..., // first two args are the same
                 void*  (*start_routine)(int),
                 int    arg);
```

- ◆ Input is anything (usually a pointer to struct for multiple arguments or even internal returns), return an integer:

```
int
pthread_create(..., // first two args are the same
                 int   (*start_routine)(void*),
                 void*  arg);
```



# Example: Creating a Thread

```
#include <pthread.h>

typedef struct_myarg_t {
    int a;
    int b;
} myarg_t;

void *mythread(void *arg) {
    myarg_t *m = (myarg_t *) arg;
    printf("%d %d\n", m->a, m->b);
    return NULL;
}

int main(int argc, char *argv[]) {
    pthread_t p;
    int rc;

    myarg_t args;
    args.a = 10;
    args.b = 20;
    rc = pthread_create(&p, NULL, mythread, &args);
    ...
}
```

# Wait for a thread to complete

```
int pthread_join(pthread_t thread, (void *)*value_ptr);
```

- ◆ `thread`: Specify which thread *to wait for*
- ◆ `value_ptr`: A pointer we want to put the return value of the start routine (ouch!)
  - Because `pthread_join()` routine changes the value, you need to **pass in a pointer** to that value.
- ◆ Returns 0 if good, or EINVAL, ESRCH if err

# Example: Waiting for Thread Completion

```
1 #include <stdio.h>
2 #include <pthread.h>
3 #include <assert.h>
4 #include <stdlib.h>
5
6 typedef struct __myarg_t {
7     int a;
8     int b;
9 } myarg_t;
10
11 typedef struct __myret_t {
12     int x;
13     int y;
14 } myret_t;
15
16 void *mythread(void *arg) {
17     myarg_t *m = (myarg_t *) arg;
18     printf("%d %d\n", m->a, m->b);
19     myret_t *r = Malloc(sizeof(myret_t));
```

```
20 r->x = 1;
21 r->y = 2;
22 return (void *) r;
23 }
24
25 int
26 main(int argc, char *argv[]) {
27 int rc;
28 pthread_t p;
29 myret_t *m;
30
31 myarg_t args;
32 args.a = 10;
33 args.b = 20;
34 Pthread_create(&p, NULL, mythread, &args);
35 Pthread_join(p, (void **) &m);
36 printf("returned %d %d\n", m->x, m->y);
37 return 0;
38 }
```

# Example: Dangerous code

- ▣ Be careful with how values are returned from a thread.

```
1  void *mythread(void *arg) {  
2  myarg_t *m = (myarg_t *) arg; 3  
    printf("%d %d\n", m->a, m->b);  
4  myret_t r; // ALLOCATED ON STACK: BAD!  
5  r.x = 1;  
6  r.y = 2;  
7  return (void *) &r;  
8  }
```

- ◆ When the variable `r` returns, it is automatically **de-allocated**.
- ◆ Don't malloc here! (memory leak prone) [bad example before]
  - Better to be consistent allocate and free in parent

# Example: Simpler Argument Passing to a Thread

- ▣ Just passing in a single value

```
1  void *mythread(void *arg) {
2      int m = (int) arg;
3      printf("%d\n", m);
4  return (void *) (arg + 1); 5
   }
6
7  int main(int argc, char *argv[]) {
8      pthread_t p;
9      int rc, m;
10     pthread_create(&p, NULL, mythread, (void *) 100);
11     pthread_join(p, (void **) &m);
12     printf("returned %d\n", m);
13 return 0; 14
   }
```

From a practical perspective  
using threads this way is pointless!  
(just do a procedure call)

- ▣ Provide **mutual exclusion** to a critical section

- ◆ Interface

```
int pthread_mutex_lock(pthread_mutex_t *mutex);  
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

- ◆ Usage (w/o *lock initialization* and *error check*)

```
pthread_mutex_t lock;  
pthread_mutex_lock(&lock);  
x = x + 1; // or whatever your critical section is  
pthread_mutex_unlock(&lock);
```

- No other thread holds the lock → the thread will acquire the lock and **enter the critical section**.
- If another thread hold the lock → the thread will **not return from the call** until it has acquired the lock.

# Locks (Cont.)

## (Cont.)

- ▣ All locks must be **properly initialized** (i.e. unlocked value).

- ◆ One way: using `PTHREAD_MUTEX_INITIALIZER`

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
```

- ◆ The dynamic way: using `pthread_mutex_init()`

```
int rc = pthread_mutex_init(&lock, NULL);  
assert(rc == 0 && "Error in mutex init");
```



# Locks (Cont.)

## (Cont.)

- ▣ Check errors code when calling lock and unlock
  - ◆ An example wrapper

```
// Use this to keep your code clean but check for failures
// Only use if exiting program is OK upon failure
void Pthread_mutex_lock(pthread_mutex_t *mutex) {
    int rc = pthread_mutex_lock(mutex);
    assert(rc == 0 && "Error in acquire");
}
```

- ▣ These two calls are used in lock acquisition

```
int pthread_mutex_trylock(pthread_mutex_t *mutex);
int pthread_mutex_timelock(pthread_mutex_t *mutex,
                           struct timespec *abs_timeout);
```

- ◆ trylock: return failure if the lock is already held
- ◆ timelock: return after a timeout

# Not a bad idea to define a wrapper: much cleaner code

```
1 #ifndef __MYTHREADS_h_
2 #define __MYTHREADS_h_
3
4 #include <pthread.h>
5 #include <assert.h>
6 #include <sched.h>
7
8 void
9 Pthread_mutex_lock(pthread_mutex_t *m)
10 {
11     int rc = pthread_mutex_lock(m);
12     assert(rc == 0);
13 }
14
15 void
16 Pthread_mutex_unlock(pthread_mutex_t *m)
17 {
18     int rc = pthread_mutex_unlock(m);
19     assert(rc == 0);
20 }
21
22 void
23 Pthread_create(pthread_t *thread, const pthread_attr_t *attr,
24               void *(*start_routine)(void*), void *arg)
25 {
26     int rc = pthread_create(thread, attr, start_routine, arg);
27     assert(rc == 0);
28 }
29
30 void
31 Pthread_join(pthread_t thread, void **value_ptr)
32 {
33     int rc = pthread_join(thread, value_ptr);
34     assert(rc == 0);
35 }
36
37 #endif // __MYTHREADS_h_
```

# Locks (Cont.)

## (Cont.)

- ▣ These two calls are also used in **lock acquisition**

```
int pthread_mutex_trylock(pthread_mutex_t *mutex);  
int pthread_mutex_timelock(pthread_mutex_t *mutex,  
                           struct timespec *abs_timeout);
```

- ◆ **trylock**: return failure if the lock is already held
- ◆ **timelock**: return after a timeout or after acquiring the lock

# Condition Variables

- ▣ **Condition variables** are useful when some kind of **signaling** must take place between threads.

```
int pthread_cond_wait(pthread_cond_t *cond,  
                      pthread_mutex_t *mutex);  
int pthread_cond_signal(pthread_cond_t *cond);
```

- ◆ `pthread_cond_wait`:
  - Put the calling thread to sleep.
  - Wait for some other thread to signal it.
- ◆ `pthread_cond_signal`:
  - Unblock at least one of the threads that are blocked on the condition variable

# Condition Variables (Cont.)

## (Cont.)

### ▣ A thread calling wait routine:

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t init = PTHREAD_COND_INITIALIZER;

pthread_mutex_lock(&lock);
while (initialized == 0)
    pthread_cond_wait(&init, &lock);
pthread_mutex_unlock(&lock);
```

- ◆ The wait call **releases the lock** when putting said caller to sleep.
- ◆ Before returning after being woken, the wait call **re-acquire the lock**.

### ▣ A thread calling signal routine:

```
pthread_mutex_lock(&lock);
initialized = 1;
pthread_cond_signal(&init);
pthread_mutex_unlock(&lock);
```

# Condition Variables (Cont.)

## (Cont.)

- ▣ The waiting thread **re-checks** the condition **in a while loop**, instead of a simple if statement.

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;  
pthread_cond_t init = PTHREAD_COND_INITIALIZER;  
  
pthread_mutex_lock(&lock);  
while (initialized == 0)  
    pthread_cond_wait(&init, &lock);  
pthread_mutex_unlock(&lock);
```

- ◆ Without rechecking, the waiting thread will continue thinking that the condition has changed *even though it has not*.
- ◆ For example if multiple threads are waiting and only one should grab the data (producer-consumer)

# Condition Variables (Cont.)

## (Cont.)

- ❑ Don't ever to this.

- ◆ A thread calling wait routine:

```
while(initialized == 0)
    ; // spin
```

- ◆ A thread calling signal routine:

```
initialized = 1;
```

- ◆ It performs poorly in many cases. → just wastes CPU cycles.
- ◆ It is error prone.

# Compiling and Running

- To compile them, you must include the header `pthread.h`
  - ◆ Explicitly link with the **pthread library**, by adding the `-pthread` flag.

```
prompt> gcc -o main main.c -Wall -pthread
```

- ◆ For more information,

```
man -k pthread
```



# Thread API Use Guidelines

- ❑ Keep it simple
  - ◆ Tricky thread interactions lead to (hard to find) bugs
- ❑ Minimize thread interaction
  - ◆ Limits scalability
- ❑ Initialize mutex and cond vars
- ❑ Check always return codes
- ❑ Be careful how to pass arguments and get values:
  - ◆ A good practice is to allocate/free memory in the calling thread
  - ◆ Be careful with heap
- ❑ Each thread has his own stack
- ❑ Always use cond. variables to signal between threads
- ❑ **Read the man pages**